Meta-date:

In AWS (Amazon Web Services), metadata refers to the information associated with AWS resources. It provides descriptive data about the resources and can include details such as configuration settings, attributes, and other relevant information.

In the context of AWS EC2 (Elastic Compute Cloud) instances, each instance has metadata that can be accessed from within the instance. This metadata includes information like the instance ID, instance type, IP addresses, security groups, and more. It can be useful for applications running on the instance to retrieve this metadata dynamically and adapt their behavior based on the environment.

The metadata can be accessed by making HTTP requests to a special URL provided by the EC2 instance itself. The URL typically follows the pattern: http://169.254.169.254/latest/meta-data/. By appending specific paths or keys to this URL, you can retrieve different types of metadata. For example, http://169.254.169.254/latest/meta-data/instance-id retrieves the instance ID.

Metadata in AWS goes beyond EC2 instances. Various AWS services have their own metadata that can be accessed in different ways. For example, AWS Identity and Access Management (IAM) roles have associated metadata that provides details about the role and its permissions.

Overall, metadata in AWS plays an important role in providing information about resources, enabling dynamic adaptation of applications, and facilitating the interaction and management of AWS services.

Metadata url:

`curl http://169.254.169.254/latest/meta-data`

For getting metadata need to enable the following option when create a ec2 instance and we can get this under advanced details section

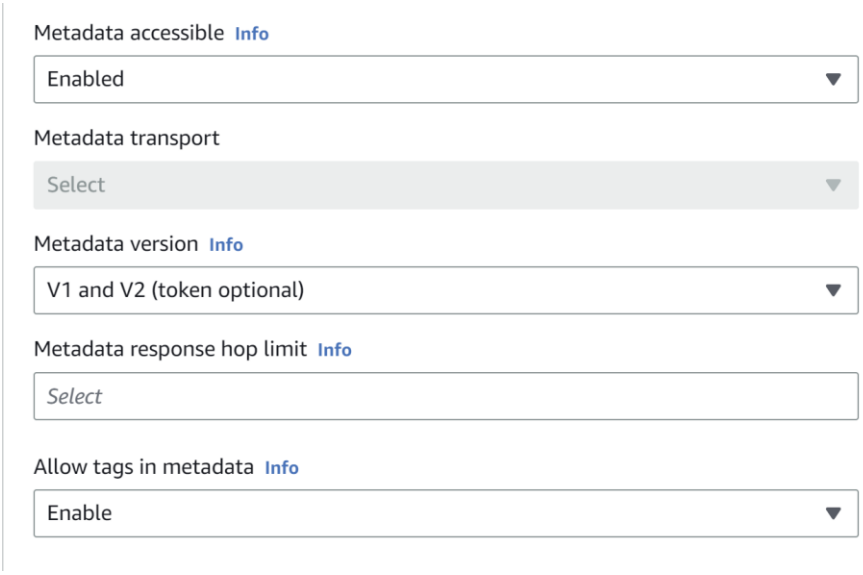

Launch Template:

An EC2 instance launch template is a configuration template that defines the settings for launching EC2 instances within an AWS account. It simplifies the process of launching instances by allowing you to create reusable templates with predefined configurations.

When you create a launch template, you can specify various parameters, including:

1. AMI (Amazon Machine Image): The template can specify the Amazon Machine Image (AMI) to use as the base for the instance.
2. Instance type: You can define the instance type, such as t2.micro, m5.large, etc., which determines the hardware specifications (CPU, memory, storage, etc.) of the EC2 instance.
3. Network settings: You can configure the networking aspects, including the VPC (Virtual Private Cloud), subnet, security groups, and IP addressing.
4. Storage options: Launch templates allow you to define the storage options for your instances, such as the root volume size, EBS (Elastic Block Store) settings, and instance store volumes.
5. Security settings: You can specify the security groups and IAM (Identity and Access Management) roles to associate with the instances launched from the template.
6. User data and startup scripts: Launch templates support adding user data, which can be used to run scripts or execute commands when the instance starts.
7. Instance tags: You can include tags to categorize and manage your instances more effectively.

By using launch templates, you can create a standardized configuration for your instances and launch them easily with consistent settings. This helps automate and streamline the process of deploying EC2 instances within your AWS infrastructure. Additionally, launch templates are compatible with other AWS services such as Auto Scaling, allowing you to scale your instances efficiently based on demand.

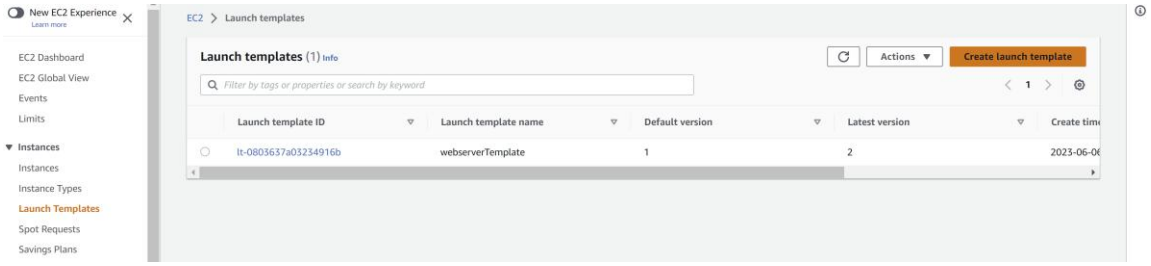

now launch an instance from that template

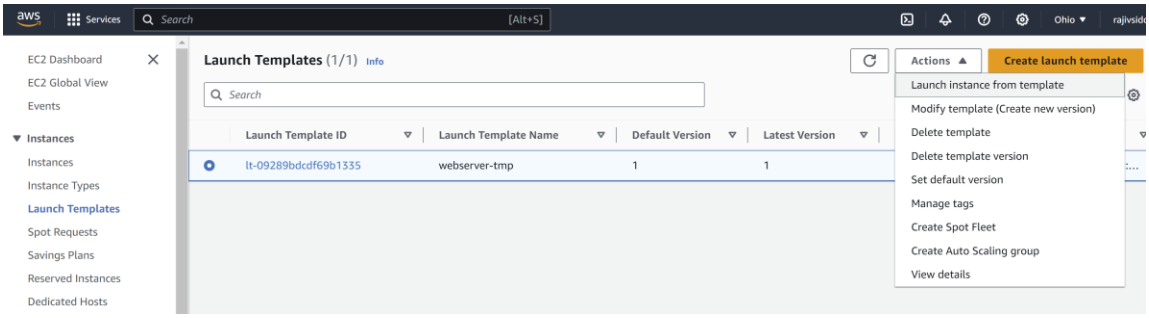

Modify the launch template then launch an instance from that template’s new version.

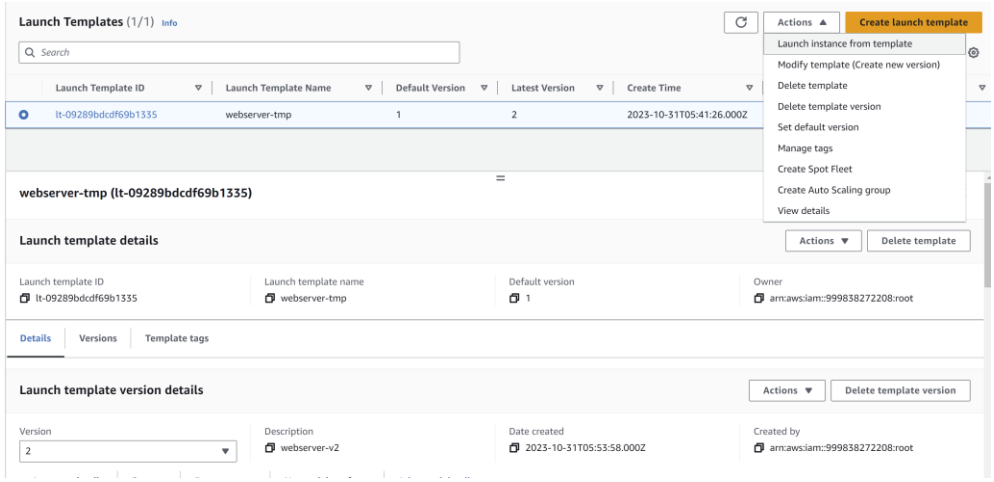

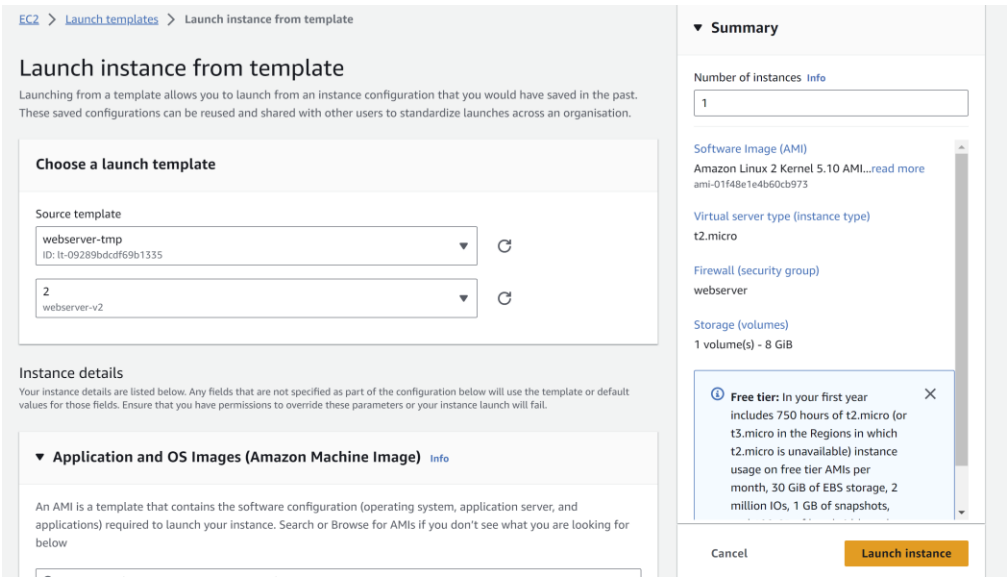

Ec2 instance purchasing options:

- On-Demand Instances – short workload, predictable pricing, pay by second
- Reserved (1 & 3 years)
  - Reserved Instances – long workloads
  - Convertible Reserved Instances – long workloads with flexible instances
- Savings Plans (1 & 3 years) –commitment to an amount of usage, long workload
- Spot Instances – short workloads, cheap, can lose instances (less reliable)
- Dedicated Hosts – book an entire physical server, control instance placement
- Dedicated Instances – no other customers will share your hardware
- Capacity Reservations – reserve capacity in a specific AZ for any duration

Ec2 On Demand

- Pay for what you use:
  - Linux or Windows - billing per second, after the first minute
  - All other operating systems - billing per hour
- Has the highest cost but no upfront payment
- No long-term commitment

- Recommended for **short-term** and **un-interrupted workloads**, where you can't predict how the application will behave

Ec2 Reserved instance

- Up to 72% discount compared to On-demand
- You reserve a specific instance attributes **(Instance Type, Region, Tenancy, OS)**
- **Reservation Period** – **1 year** (+discount) or **3 years** (+++discount)
- **Payment Options** – **No Upfront** (+), **Partial Upfront** (++), **All Upfront** (+++)
- **Reserved Instance's Scope** – **Regional** or **Zonal** (reserve capacity in an AZ)
- Recommended for steady-state usage applications (think database)
- You can buy and sell in the Reserved Instance Marketplace

- Convertible Reserved Instance
  - Can change the EC2 instance type, instance family, OS, scope and tenancy
  - Up to 66% discount

Ec2 saving plane

- Get a discount based on long-term usage (up to 72% - same as RIs)
- Commit to a certain type of usage ($10/hour for 1 or 3 years)
- Usage beyond EC2 Savings Plans is billed at the On-Demand price

- Locked to a specific instance family & AWS region (e.g., M5 in us-east-1)
- Flexible across:
  - Instance Size (e.g., m5.xlarge, m5.2xlarge)
  - OS (e.g., Linux, Windows)
  - Tenancy (Host, Dedicated, Default)

Ec2 spot instance

- Can get a **discount of up to 90%** compared to On-demand
- Instances that you can "lose" at any point of time if your max price is less than the current spot price
- The **MOST cost-efficient** instances in AWS

- **Useful for workloads that are resilient to failure**
  - Batch jobs
  - Data analysis
  - Image processing
  - Any **distributed** workloads
  - Workloads with a flexible start and end time

- **Not suitable for critical jobs or databases**

Ec2 dedicated host:

- A physical server with EC2 instance capacity fully dedicated to your use
- Allows you address **compliance requirements** and **use your existing server-bound software licenses** (per-socket, per-core, pe—VM software licenses)

- Purchasing Options:
  - **On-demand** – pay per second for active Dedicated Host
  - **Reserved** - 1 or 3 years (No Upfront, Partial Upfront, All Upfront)
- The most expensive option

- Useful for software that have complicated licensing model (BYOL – Bring Your Own License)
- Or for companies that have strong regulatory or compliance needs

Ec2 Dedicated instances:

- Instances run on hardware that's dedicated to you

- May share hardware with other instances in same account

- No control over instance placement (can move hardware after Stop / Start)

Characteristic	Dedicated Instances	Dedicated Hosts
Enables the use of dedicated physical servers	X	X
Per instance billing (subject to a $2 per region fee)	X	
Per host billing		X
Visibility of sockets, cores, host ID		X
Affinity between a host and instance		X
Targeted instance placement		X
Automatic instance placement	X	X
Add capacity using an allocation request		X

Ec2 capacity reservation

- Reserve **On-Demand** instances capacity in a specific AZ for any duration
- You always have access to EC2 capacity when you need it
- **No time commitment** (create/cancel anytime), **no billing discounts**
- Combine with Regional Reserved Instances and Savings Plans to benefit from billing discounts
- You're charged at On-Demand rate whether you run instances or not

- Suitable for short-term, uninterrupted workloads that needs to be in a specific AZ

Explain with a practical example:

- On demand: coming and staying in resort whenever we like, we pay the full price
- Reserved: like planning ahead and if we plan to stay for a long time, we may get a good discount.
- Savings Plans: pay a certain amount per hour for certain period and stay in any room type (e.g., King, Suite, Sea View, …)
- Spot instances: the hotel allows people to bid for the empty rooms and the highest bidder keeps the rooms. You can get kicked out at any time
- Dedicated Hosts: We book an entire building of the resort
- Capacity Reservations: you book a room for a period with full price even you don't stay in it

Price comparison example – m4.large -us east 1

Price Type	Price (per hour)
On-Demand	$0.10
Spot Instance (Spot Price)	$0.038 - $0.039 (up to 61% off)
Reserved Instance (1 year)	$0.062 (No Upfront) - $0.058 (All Upfront)
Reserved Instance (3 years)	$0.043 (No Upfront) - $0.037 (All Upfront)
EC2 Savings Plan (1 year)	$0.062 (No Upfront) - $0.058 (All Upfront)
Reserved **Convertible** Instance (1 year)	$0.071 (No Upfront) - $0.066 (All Upfront)
Dedicated Host	On-Demand Price
Dedicated Host Reservation	Up to 70% off
Capacity Reservations	On-Demand Price

Spot instance

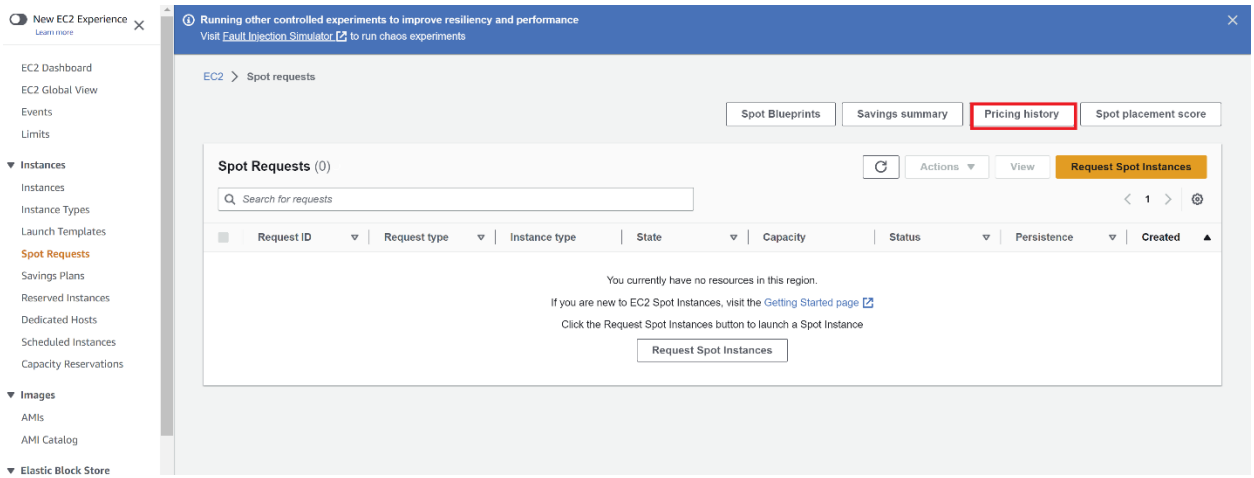

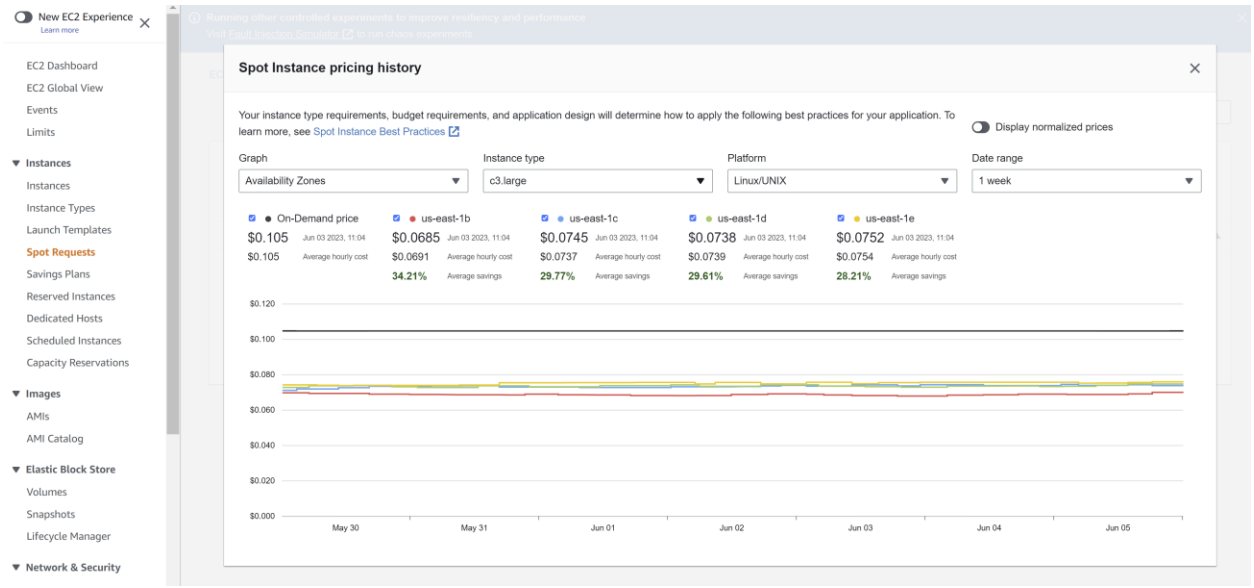

Amazon Machine Image(AMI):

In the context of Amazon Web Services (AWS), an AMI stands for Amazon Machine Image. An AMI is a pre-configured template that contains the necessary information to launch an EC2 instance. It includes the operating system, software packages, configuration settings, and any additional data required for the instance.

Here are some key points about AMIs in AWS:

1. Pre-built Templates: AMIs provide a starting point for launching instances. AWS offers a wide range of pre-built AMIs that are tailored for various use cases, including different operating systems (such as Amazon Linux, Ubuntu, Windows Server, etc.) and application-specific configurations.

2. Public and Private AMIs: AWS provides a public repository of AMIs that are available to all AWS users. These AMIs are created and shared by the AWS community or AWS itself. Additionally, you can create your own private AMIs by customizing and saving an existing EC2 instance as an image.

3. Customization and Configuration: You can customize an AMI to meet your specific requirements. This involves installing additional software, applying updates, configuring security settings, and making any desired changes to the base image. Once the modifications are made, you can save the customized instance as a new AMI.

4. Versioning and Lifecycle: AMIs support versioning, allowing you to manage and track different iterations of an image. You can create new versions of an AMI as you make updates or changes to the underlying configuration. Additionally, you can control the lifecycle of an AMI, including when to deprecate or deregister an image.

5. Marketplace AMIs: AWS Marketplace offers a wide selection of AMIs from third-party vendors. These AMIs may come with pre-installed software, databases, or complete application stacks, enabling you to quickly deploy complex solutions without having to manually configure them.

To launch an EC2 instance, you typically select an appropriate AMI, specify the instance type, configure networking, security settings, and storage options. By leveraging AMIs, you can easily replicate and deploy instances with the desired operating system and software configurations.

Elastic ip:

Add a new elastic ip

Ec2>Network & security>Elastic Ips>click allocate elastic ip button>click allocate

Attach this elastic ip to a ec2 instance

Ec2>Network & security>Elastic Ips>select the elastic ip>click Action >click Associate Elastic ip

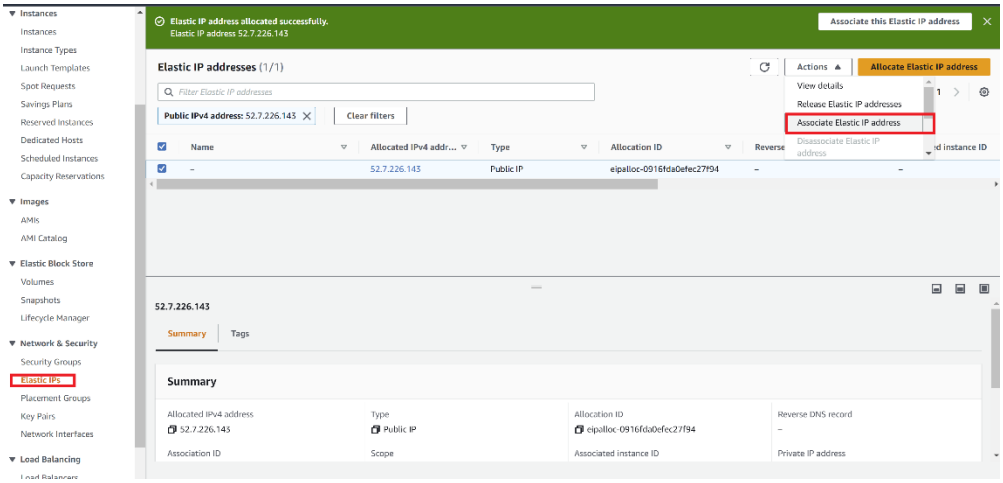

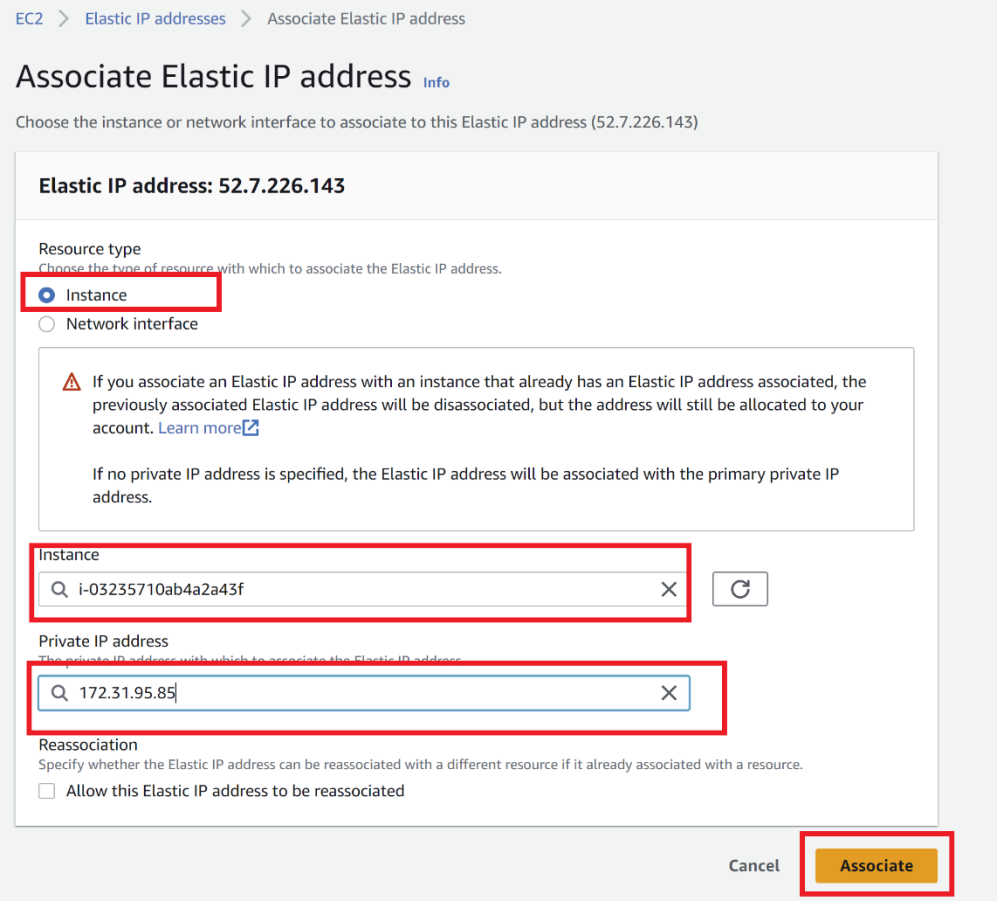

Placement Groups:

- Sometimes you want control over the EC2 Instance placement strategy
- That strategy can be defined using placement groups
- When you create a placement group, you specify one of the following strategies for the group:
  - *Cluster*—clusters instances into a low-latency group in a single Availability Zone
  - *Spread*—spreads instances across underlying hardware (max 7 instances per group per AZ)
  - *Partition*—spreads instances across many different partitions (which rely on different sets of racks) within an AZ. Scales to 100s of EC2 instances per group (Hadoop, Cassandra, Kafka)

Placement Groups-Cluster

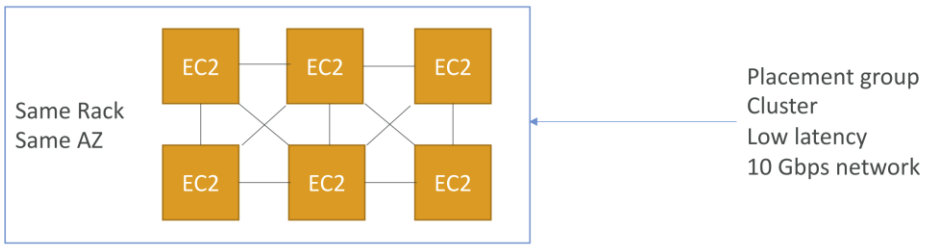

- Pros: Great network (10 Gbps bandwidth between instances with Enhanced Networking enabled - recommended)
- Cons: If the rack fails, all instances fails at the same time
- Use case:
  - Big Data job that needs to complete fast
  - Application that needs extremely low latency and high network throughput

Placement Groups- Spread

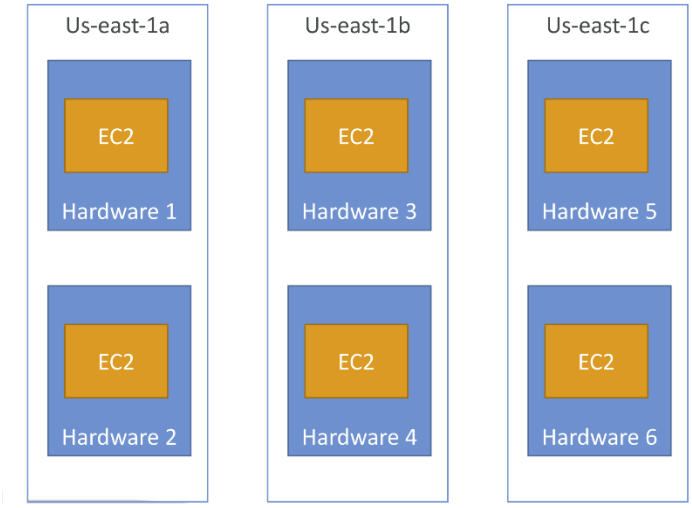

- Pros:
  - Can span across Availability Zones (AZ)
  - Reduced risk is simultaneous failure
  - EC2 Instances are on different physical hardware
- Cons:
  - Limited to 7 instances per AZ per placement group
- Use case:
  - Application that needs to maximize high availability
  - Critical Applications where each instance must be isolated from failure from each other

Placement Groups- Partition

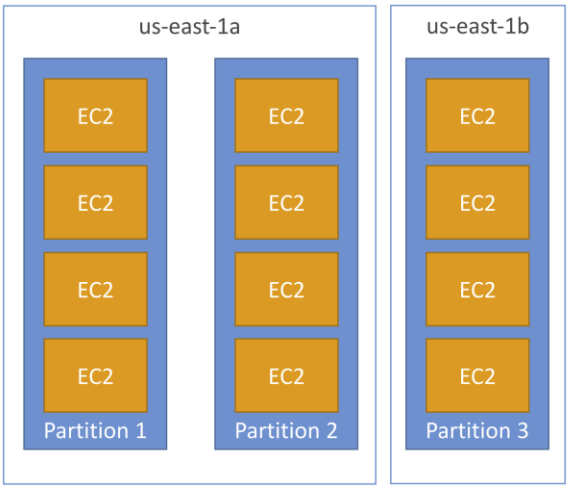

- Up to 7 partitions per AZ
- Can span across multiple AZs in the same region
- Up to 100s of EC2 instances
- The instances in a partition do not share racks with the instances in the other partitions
- A partition failure can affect many EC2 but won't affect other partitions
- EC2 instances get access to the partition information as metadata
- Use cases: HDFS, HBase, Cassandra, Kafka

Ec2 hibernate:

- Introducing EC2 Hibernate:
  - The in-memory (RAM) state is preserved
  - The instance boot is much faster! (the OS is not stopped / restarted)
  - Under the hood: the RAM state is written to a file in the root EBS volume
  - The root EBS volume must be encrypted
- Use cases:
  - Long-running processing
  - Saving the RAM state
  - Services that take time to initialize

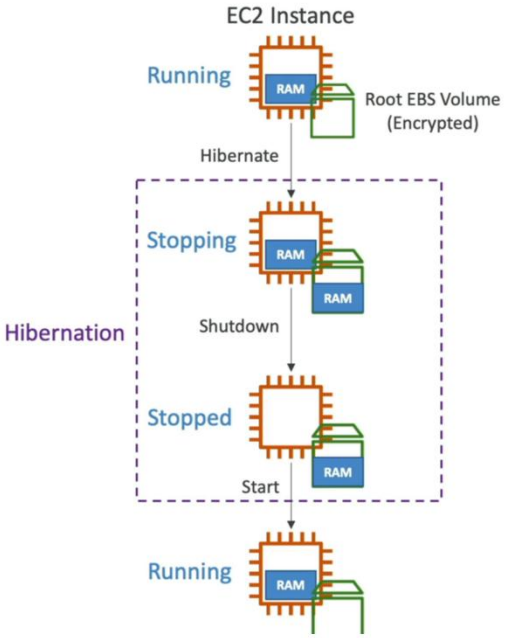

Hibernate- Good to know

- Supported Instance Families – C3, C4, C5, I3, M3, M4, R3, R4, T2, T3, …
- Instance RAM Size – must be less than 150 GB.
- Instance Size – not supported for bare metal instances.
- AMI – Amazon Linux 2, Linux AMI, Ubuntu, RHEL, CentOS & Windows…
- Root Volume – must be EBS, encrypted, not instance store, and large
- Available for On-Demand, Reserved and Spot Instances

- An instance can NOT be hibernated more than 60 days

Enable hibernate when launching an ec2 instance and follow showing the steps.
ec2>launch instance> scroll down and advanced details

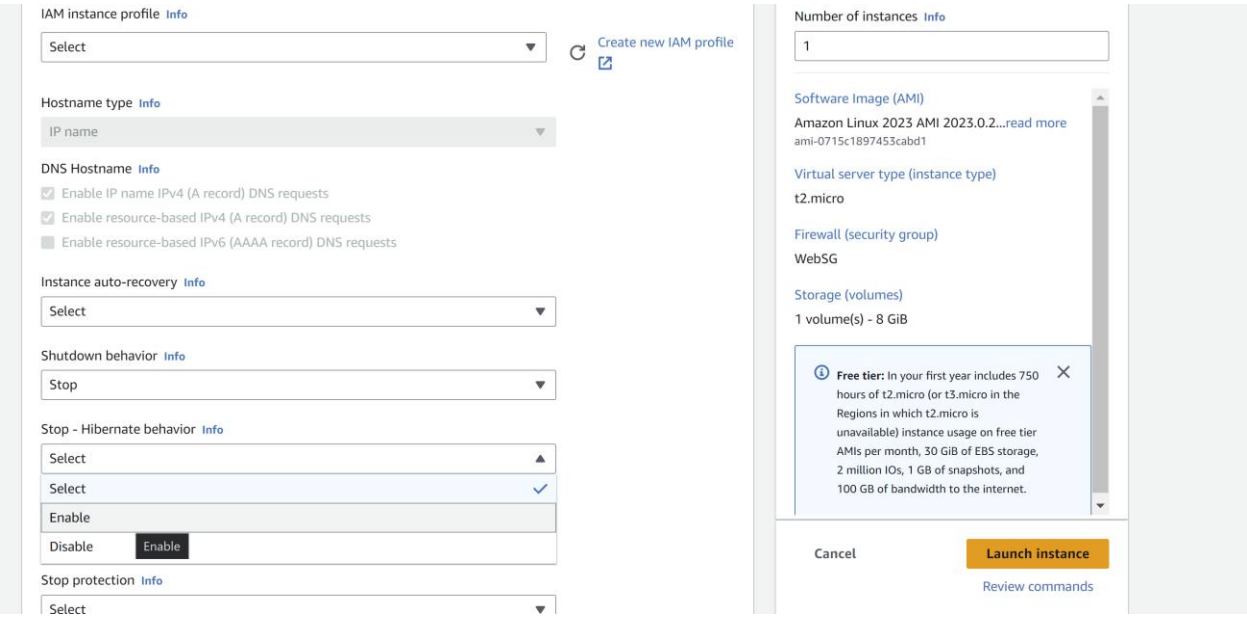

To enable hibernation, space is allocated on the root volume to store the instance memory (RAM). Make sure that the **root volume is large enough to store** the RAM contents and accommodate your expected usage, e.g. OS, applications. To use hibernation, the **root volume must be an encrypted** EBS volume.

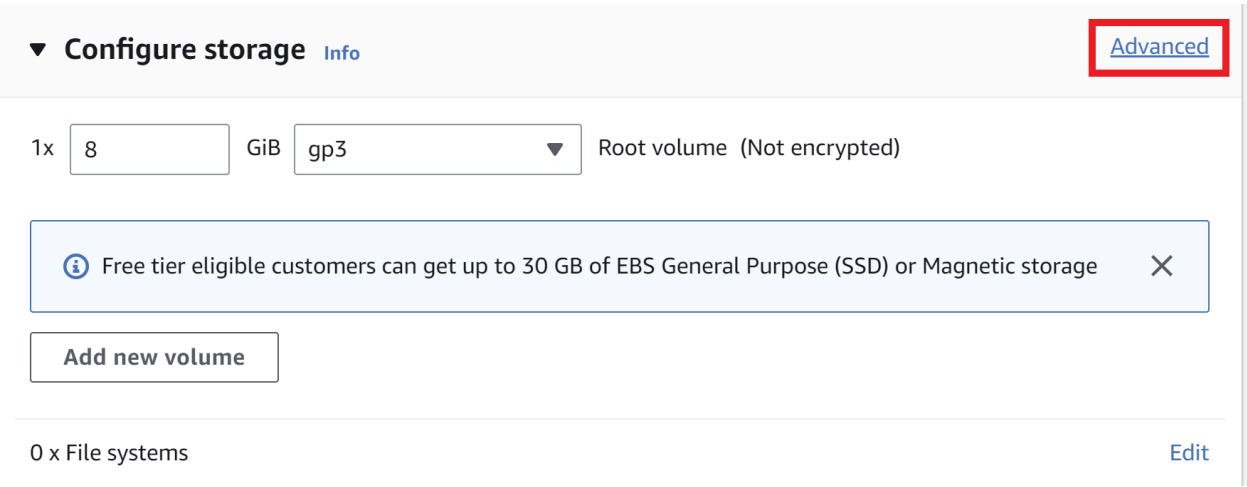

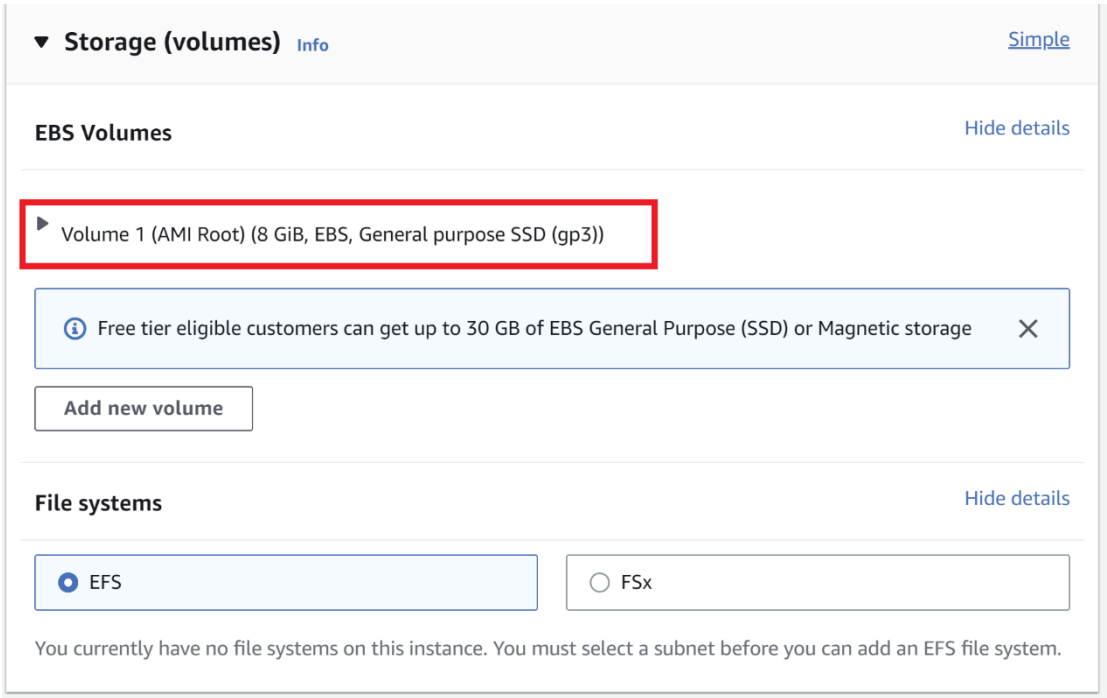
Storage (volumes) Info
Simple
EBS Volumes
Hide details
Volume 1 (AMI Root) (8 GiB, EBS, General purpose SSD (gp3))
Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage
Add new volume
File systems
Hide details
EFS
FSx
You currently have no file systems on this instance. You must select a subnet before you can add an EFS file system.

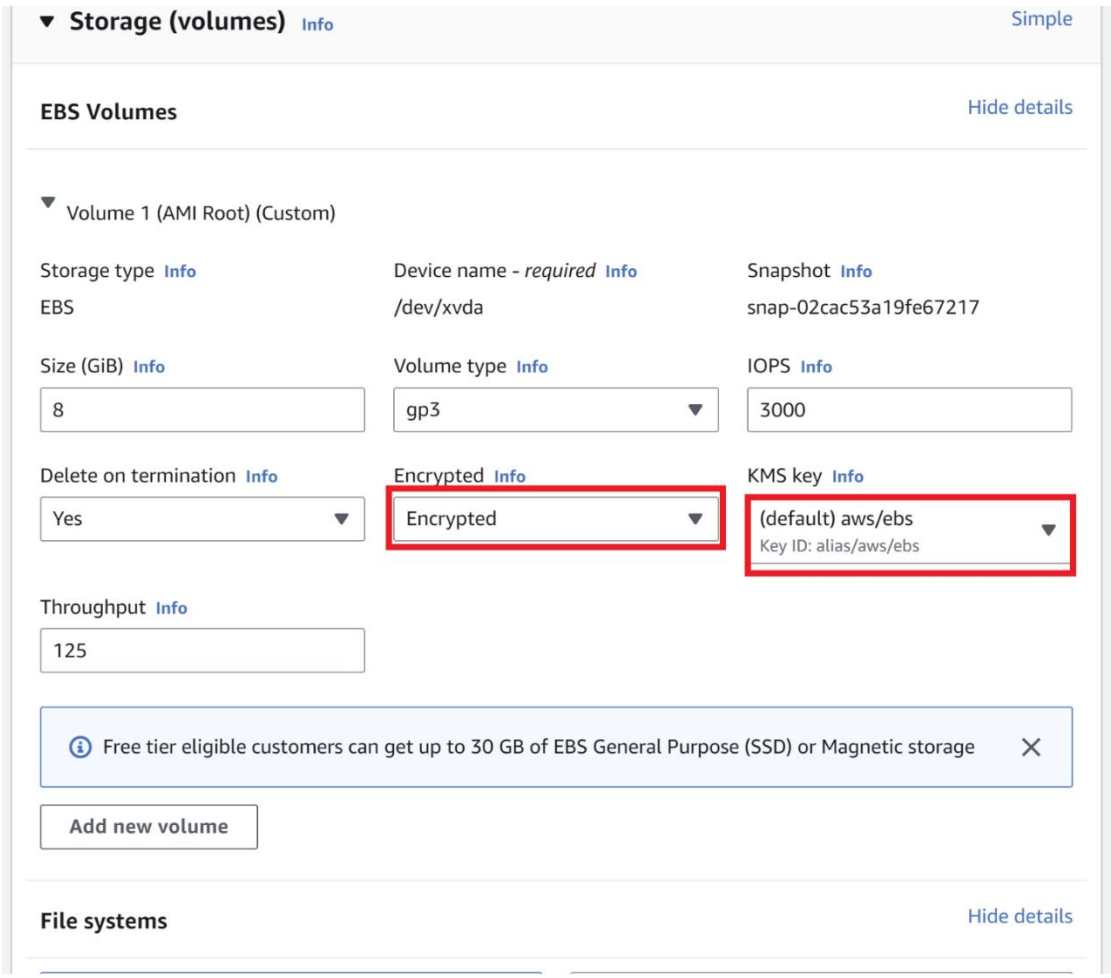
Storage (volumes) Info
Simple
EBS Volumes
Hide details
Volume 1 (AMI Root) (Custom)
Storage type Info
EBS
Device name - required Info
/dev/xvda
Snapshot Info
snap-02cac53a19fe67217
Size (GiB) Info
8
Volume type Info
gp3
IOPS Info
3000
Delete on termination Info
Yes
Encrypted Info
Encrypted
KMS key Info
(default) aws/ebs
Key ID: alias/aws/ebs
Throughput Info
125
Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage
Add new volume
File systems
Hide details

Now launch the instance and login and type the following command for checking the system uptime
#uptime

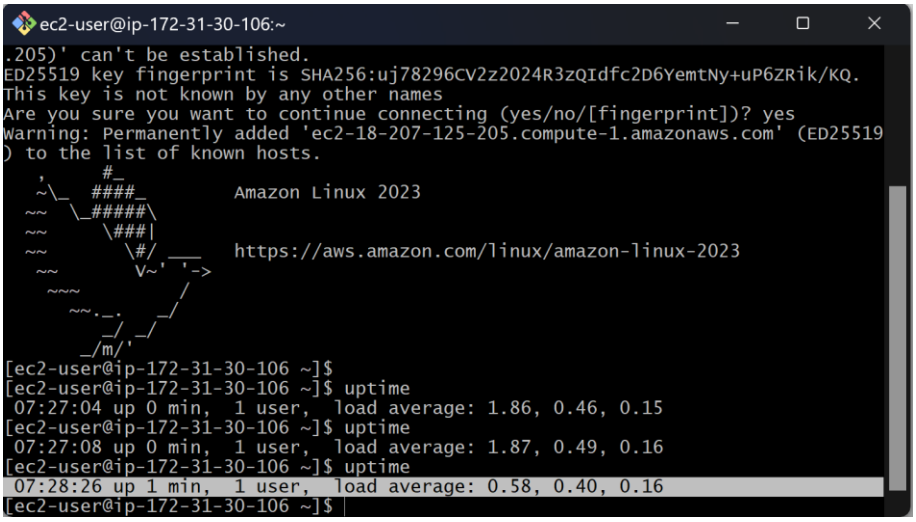

Now hibernate the instance as shown the following picture

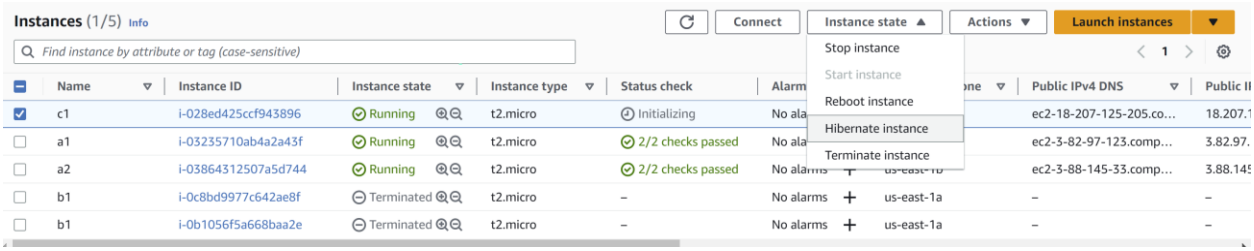

Now start the instance and login and check the uptime and you can see it will showing more uptime